

```

from google.colab import drive
drive.mount('/content/drive')

Mounted at /content/drive

import os
import time
import torch
import torch.nn as nn
import torchvision
from torchvision import transforms

import numpy as np
from tqdm import tqdm
import matplotlib.pyplot as plt

transform = transforms.ToTensor()

train_dataset = torchvision.datasets.MNIST(root = "./data" , train = True , download = True , transform = transform)
valid_dataset = torchvision.datasets.MNIST(root = "./data" , train = False , download = True , transform = transform)

Downloading http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz
Downloading http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz to ./data/MNIST/raw/train-images-idx3-ubyte.gz
100%|██████████| 9912422/9912422 [00:00<00:00, 143944767.85it/s]
Extracting ./data/MNIST/raw/train-images-idx3-ubyte.gz to ./data/MNIST/raw

Downloading http://yann.lecun.com/exdb/mnist/train-labels-idx1-ubyte.gz
Downloading http://yann.lecun.com/exdb/mnist/train-labels-idx1-ubyte.gz to ./data/MNIST/raw/train-labels-idx1-ubyte.gz
100%|██████████| 28881/28881 [00:00<00:00, 47541481.09it/s]
Extracting ./data/MNIST/raw/train-labels-idx1-ubyte.gz to ./data/MNIST/raw

Downloading http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz
Downloading http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz to ./data/MNIST/raw/t10k-images-idx3-ubyte.gz
100%|██████████| 1648877/1648877 [00:00<00:00, 50785300.20it/s]Extracting ./data/MNIST/raw/t10k-images-idx3-ubyte.gz to ./data/MNIST/raw

Downloading http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz
Downloading http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz to ./data/MNIST/raw/t10k-labels-idx1-ubyte.gz
100%|██████████| 4542/4542 [00:00<00:00, 7186167.02it/s]
Extracting ./data/MNIST/raw/t10k-labels-idx1-ubyte.gz to ./data/MNIST/raw

train_dl = torch.utils.data.DataLoader(train_dataset , batch_size = 100)

class Encoder(nn.Module):
    def __init__(self , input_size = 28*28 , hidden_size1 = 500, hidden_size2 = 250 , hidden_size3 = 100, hidden_size4 = 50, z_dim
        super().__init__()
        self.fc1 = nn.Linear(input_size , hidden_size1)
        self.fc2 = nn.Linear(hidden_size1 , hidden_size2)
        self.fc3 = nn.Linear(hidden_size2 , hidden_size3)
        self.fc4 = nn.Linear(hidden_size3 , hidden_size4)
        self.fc5 = nn.Linear(hidden_size4 , z_dim)
        self.relu = nn.ReLU()
    def forward(self , x):
        x = self.relu(self.fc1(x))
        x = self.relu(self.fc2(x))
        x = self.relu(self.fc3(x))
        x = self.relu(self.fc4(x))
        x = self.fc5(x)
        return x

class Decoder(nn.Module):
    def __init__(self , output_size = 28*28 , hidden_size1 = 500 , hidden_size2 = 250 , hidden_size3 = 100, hidden_size4 = 50, z_di
        super().__init__()
        self.fc1 = nn.Linear(z_dim , hidden_size4)
        self.fc2 = nn.Linear(hidden_size4 , hidden_size3)
        self.fc3 = nn.Linear(hidden_size3 , hidden_size2)
        self.fc4 = nn.Linear(hidden_size2 , hidden_size1)
        self.fc5 = nn.Linear(hidden_size1 , output_size)
        self.relu = nn.ReLU()
    def forward(self , x):
        x = self.relu(self.fc1(x))

```

```

    x = self.relu(self.fc2(x))
    x = self.relu(self.fc3(x))
    x = self.relu(self.fc4(x))
    x = torch.sigmoid(self.fc5(x))
    return x

device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
device

device(type='cuda')

enc = Encoder().to(device)
dec = Decoder().to(device)

loss_fn = nn.MSELoss()
optimizer_enc = torch.optim.Adam(enc.parameters())
optimizer_dec = torch.optim.Adam(dec.parameters())

train_loss = []
num_epochs = 400
checkpoint_path = '/content/drive/MyDrive/model/Autoencoder/checkpoint_30z_5h_200e.pth'

# Check if a checkpoint exists to resume training
if os.path.exists("/content/drive/MyDrive/model/Autoencoder/checkpoint_30z_5h_200e.pth"):
    checkpoint = torch.load("/content/drive/MyDrive/model/Autoencoder/checkpoint_30z_5h_200e.pth")
    enc.load_state_dict(checkpoint["enc_state_dict"])
    dec.load_state_dict(checkpoint["dec_state_dict"])
    optimizer_enc.load_state_dict(checkpoint["optimizer_enc_state_dict"])
    optimizer_dec.load_state_dict(checkpoint["optimizer_dec_state_dict"])
    train_loss = checkpoint["loss"]
    start_epoch = checkpoint["epoch"] + 1 # Start from the next epoch after the loaded checkpoint
    print("Resume training from epoch", start_epoch)
else:
    start_epoch = 1

    Resume training from epoch 401

total_batches = len(train_dl)
for epoch in range(start_epoch, num_epochs+1):
    train_epoch_loss = 0
    start_time = time.time()
    # Create a tqdm progress bar for the epoch
    epoch_progress = tqdm(enumerate(train_dl, 1), total=total_batches, desc=f'Epoch {epoch}/{num_epochs}', leave=False)
    for step, (imgs, _) in epoch_progress:
        imgs = imgs.to(device)
        imgs = imgs.flatten(1)
        latents = enc(imgs)
        output = dec(latents)
        loss = loss_fn(output, imgs)
        train_epoch_loss += loss.item()
        optimizer_enc.zero_grad()
        optimizer_dec.zero_grad()
        loss.backward()
        optimizer_enc.step()
        optimizer_dec.step()
        # Update the progress bar description with current step and loss
        epoch_progress.set_description(f'Epoch {epoch}/{num_epochs}, Step {step}/{total_batches}, Loss: {loss.item():.4f}')
    train_loss.append(train_epoch_loss)
    # Close the tqdm progress bar for the epoch
    epoch_progress.close()

    # Print the epoch loss after each epoch
    print('\n')
    print(f'Epoch {epoch}/{num_epochs}, Loss: {train_epoch_loss:.4f}, Time taken: [{time.time() - start_time:.2f}s]')

    # Save the model checkpoint along with training-related information
    checkpoint = {
        'epoch': epoch,
        'enc_state_dict': enc.state_dict(), # Save the encoder model's state dictionary
        'dec_state_dict': dec.state_dict(),
        'optimizer_enc_state_dict': optimizer_enc.state_dict(), # Save the optimizer state
        'optimizer_dec_state_dict': optimizer_dec.state_dict(),
        'loss': train_loss, # Save the loss

```


```

    }
    torch.save(checkpoint, checkpoint_path)

checkpoint = torch.load("/content/drive/MyDrive/model/Autoencoder/checkpoint_30z_5h_400e.pth")
saved_losses = checkpoint['loss']

# Plot the loss values
plt.plot(saved_losses)
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.title('Training Loss')
plt.show()

```



```

values = None
all_labels = []

with torch.no_grad():
    for (imgs , labels) in train_dl:
        imgs = imgs.to(device)
        imgs = imgs.flatten(1)
        all_labels.extend(list(labels.numpy()))
        latents = enc(imgs)
        if values is None:
            values = latents.cpu()
        else:
            values = torch.vstack([values , latents.cpu()])

values.shape

torch.Size([60000, 30])

len(all_labels)

60000

# cmap = plt.get_cmap('viridis', 10)
# cmap
all_labels = np.array(all_labels)
values = values.numpy()
# pc = plt.scatter(values[:, 0] , values[:, 1] , c = all_labels , cmap = cmap)
# plt.colorbar(pc)

all_means = {}
for i in range(10):
    inds = np.argwhere(all_labels == i)
    print(inds.shape)
    num_latents = values[inds].squeeze()

```

```
mean = num_latents.mean(axis = 0)
# all_means[i] = (mean[0] , mean[1], mean[2], mean[3], mean[4])
all_means[i] = tuple(mean[:30])
```

```
(5923, 1)
(6742, 1)
(5958, 1)
(6131, 1)
(5842, 1)
(5421, 1)
(5918, 1)
(6265, 1)
(5851, 1)
(5949, 1)
```

```
all_means
```

```
{0: (-0.019664943,
-0.04140795,
-0.093276344,
-0.034774188,
0.06456965,
0.6099748,
4.649796,
0.28887,
2.0323434,
-0.006129101,
0.44975147,
-0.09560164,
0.14724708,
0.70593846,
1.0216453,
1.5119395,
-1.3156971,
0.11654582,
0.5168116,
-0.5505509,
-0.6365081,
0.010190381,
1.6171646,
0.0033239129,
-0.15580587,
1.4233271,
-0.023702528,
-0.044816338,
-0.12257131,
0.010026846),
1: (0.041038856,
-0.31787196,
-0.022012569,
0.045954015,
0.10092047,
0.48969728,
-1.9134811,
-1.8123839,
-1.8941734,
-0.02852272,
0.029468417,
-0.21950857,
0.07557216,
-0.9434995,
3.6907852,
0.31471026,
0.10005181,
-0.37161803,
0.36695603,
-0.2910862,
2.2215939,
0.0038460789,
-2.867131,
0.038801245,
-0.11711168,
-1.6802964,
1.0380716,
0.011429321,
~ .....
```

```
with torch.no_grad():
    pred = dec(torch.Tensor(all_means[5])[None , ...].to(device)).cpu()
    transforms.ToPILImage()(pred.reshape(1 , 28 , 28))
```

